



Image by Alan Warburton / © BBC / Better Images of AI / Plant / CC-BY 4.0  
<https://betterimagesofai.org/images?dimage=2>

# COMP47590: ADVANCED MACHINE LEARNING

## SUPERVISED LEARNING - EVALUATION

Dr. Brian Mac Namee



1

## Section Outline

In this section we will cover:

- Why evaluate?
- Choosing appropriate evaluation metrics
- Running machine learning experiments

This lecture largely focuses on supervised machine learning.

2

# Why Evaluate?

3


## Prosper-René Blondlot Discoverer of n-Rays



[https://en.wikipedia.org/wiki/Prosper-Ren%C3%A9\\_Blondlot](https://en.wikipedia.org/wiki/Prosper-Ren%C3%A9_Blondlot)

Who has  
heard of  
n-rays?

4



Fundamentals of Machine Learning for Predictive Data Analytics: Algorithms, Worked Examples, and Case Studies ([www.machinlearningbook.com](http://www.machinlearningbook.com))  
John D. Kelleher, Brian Mac Namee and Aoife D'Arcy

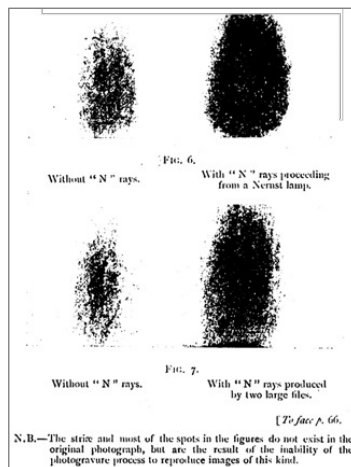


5

## Prosper-René Blondlot Discoverer of n-Rays



[https://en.wikipedia.org/wiki/Prosper-Ren%C3%A9\\_Blondlot](https://en.wikipedia.org/wiki/Prosper-Ren%C3%A9_Blondlot)



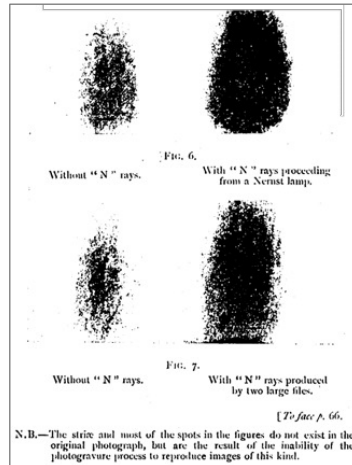
<https://en.wikipedia.org/wiki/N-ray>

6


## Prosper-René Blondlot Discoverer of n-Rays



[https://en.wikipedia.org/wiki/Prosper-Ren%C3%A9\\_Blondlot](https://en.wikipedia.org/wiki/Prosper-Ren%C3%A9_Blondlot)



<https://en.wikipedia.org/wiki/N-ray>



<https://www.wired.com/2014/09/fantastically-wrong-n-rays/>

APS This Month in Physics History:  
September 1904: Robert Wood debunks N-rays  
<https://www.aps.org/publications/apsnews/200708/history.cfm>

Fantastically Wrong: The Imaginary Radiation That Shocked Science  
and Ruined Its 'Discoverer'  
<https://www.wired.com/2014/09/fantastically-wrong-n-rays/>

7

## Why Evaluate?

It is worth distinguishing between two different types of evaluation that we do in machine learning

1. evaluating a model that we would like to deploy for a specific task (**industry**)
2. comparing machine learning methods (**research**)

8

## Why Evaluate?

The purpose of evaluation when we want to deploy a model (**industry**) is threefold

1. to determine which algorithm (plus hyperparameter values) is the most suitable for a task
2. to estimate how a model will perform after deployment
3. to convince users that a model will meet their needs

9

## Why Evaluate?

The purpose of evaluation when we want to comparing machine learning methods (**research**) can be different

1. to evaluate the performance of a new method against existing baselines
2. to determine the *best* machine learning approach for a specific problem
3. to perform a benchmark experiment

These almost all reduce to an experiment in which we compare multiple approaches using multiple datasets

10

# Evaluation For Deployment

11

## Hold-Out Test Set

It is important to establish a hold-out test set early on (we'll come back to why)



Fundamentals of Machine Learning for Predictive Data Analytics: Algorithms, Worked Examples, and Case Studies ([www.machinelearningbook.com](http://www.machinelearningbook.com))  
John D. Kelleher, Brian Mac Namee and Aoife D'Arcy



12

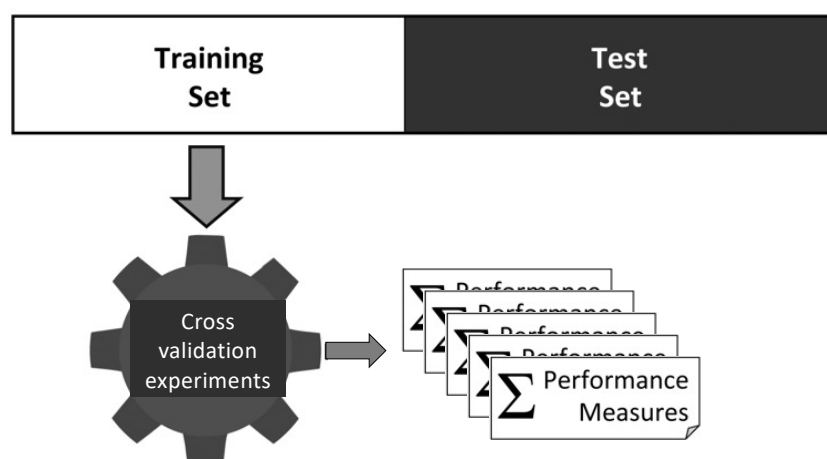
## Evaluation For Deployment

The purpose of evaluation when we want to deploy a model (**industry**) is threefold

1. **to determine which algorithm (plus hyperparameter values) is the most suitable for a task**
2. to estimate how the model will perform after deployment
3. to convince users that the model will meet their needs

13

## Evaluation For Deployment




Fundamentals of Machine Learning for Predictive Data Analytics: Algorithms,  
Worked Examples, and Case Studies ([www.machinelearningbook.com](http://www.machinelearningbook.com))  
John D. Kelleher, Brian Mac Namee and Aoife D'Arcy



14

## Evaluation For Deployment

### $k$ fold cross validation




Fundamentals of Machine Learning for Predictive Data Analytics: Algorithms, Worked Examples, and Case Studies ([www.machinlearningbook.com](http://www.machinlearningbook.com))  
John D. Kelleher, Brian Mac Namee and Aoife D'Arcy

15

## Evaluation For Deployment

### Hyper-parameters

- Hyper-parameters are parameters of the learning algorithm
- All machine learning algorithms have them!
- Hyperparameter values can make a big difference!



Model Evaluation, Model Selection, and Algorithm Selection in Machine Learning, Sebastian Raschka  
<https://arxiv.org/abs/1811.12808>

16



## Evaluation For Deployment

### Hyper-parameter tuning

- A grid search or random search based on cross validations makes sense
- AutoML approaches are good if you have access to them
- After finding best hyper-parameters we should do one last cross validation experiment with a different shuffle

17

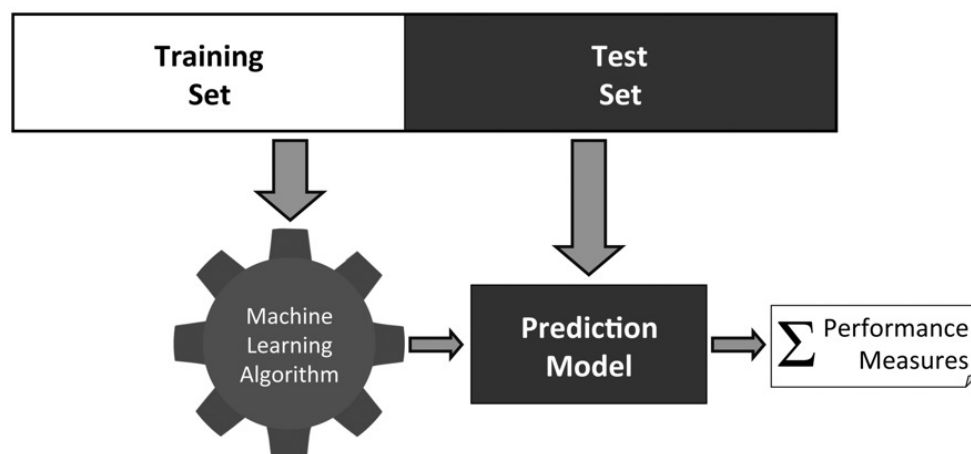
## Evaluation For Deployment

The purpose of evaluation when we want to deploy a model (**industry**) is threefold

1. to determine which model is the most suitable for a task
- 2. to estimate how the model will perform after deployment**
3. to convince users that the model will meet their needs

18

## Evaluation For Deployment



Fundamentals of Machine Learning for Predictive Data Analytics: Algorithms, Worked Examples, and Case Studies ([www.machinlearningbook.com](http://www.machinlearningbook.com))  
John D. Kelleher, Brian Mac Namee and Aoife D'Arcy



19

## Evaluation For Deployment

The purpose of evaluation when we want to deploy a model (**industry**) is threefold

1. to determine which model is the most suitable for a task
2. to estimate how the model will perform after deployment
3. **to convince users that the model will meet their needs**

20

# Performance Measures

21

## Performance Measures

Machine learning offers us a raft of different performance measures for experiments

- Accuracy
- Precision, Recall, F1 score
- Macro-averaged F1 score
- Sensitivity, Specificity
- ROC index
- Dunne index
- Gain & Lift
- Kolmogorov smirnov
- RMSE
- $R^2$
- ...

It is really important to pick the right performance measure for the problem you are solving

22

## Performance Measures

Model 1 confusion matrix **91%**

		Prediction	
		'non-churn'	'churn'
Target	'non-churn'	90	0
	'churn'	9	1

Model 2 confusion matrix **78%**

		Prediction	
		'non-churn'	'churn'
Target	'non-churn'	70	20
	'churn'	2	8

23

## Performance Measures

Model 1 confusion matrix  $\frac{1}{2}(1 + 0.1) = 55\%$

		Prediction	
		'non-churn'	'churn'
Target	'non-churn'	90	0
	'churn'	9	1

Model 2 confusion matrix  $\frac{1}{2}(0.778 + 0.8) = 78.889\%$

		Prediction	
		'non-churn'	'churn'
Target	'non-churn'	70	20
	'churn'	2	8

24

## **Macro Averaging**

The specific lesson in this example is to use macro averaging rather than micro averaging when classification datasets are imbalanced

25

## **Evaluation For Research**

26

## Why Evaluate?

The purpose of evaluation when we want to comparing machine learning methods (**research**) can be different

1. to evaluate the performance of a new method against existing baselines
2. to determine the *best* machine learning approach for a specific problem
3. to perform a benchmark experiment

These almost all reduce to an experiment in which we compare multiple approaches using multiple datasets

27

## Why Evaluate?

The purpose of evaluation when we want to comparing machine learning methods (**research**) can be different

1. to evaluate the performance of a new method against existing baselines
2. to determine the *best* machine learning approach for a specific problem
3. to perform a benchmark experiment

**We don't care about building one specific model for a problem**

These almost all reduce to an experiment in which we compare multiple approaches using multiple datasets

28

Benchmarks										
	Data 1	Data 2	Data 3	Data 4	Data 5	Data 6	Data 7	Data 8	Data 9	Data 10
Approach 1										
Approach 2										
Approach 3										
Approach 4										
Approach 5										
Approach 6										
Approach 7										
Approach 8										
Approach 9										
Approach 10										

29

Benchmarks - Performance Measures										
	Data 1	Data 2	Data 3	Data 4	Data 5	Data 6	Data 7	Data 8	Data 9	Data 10
Approach 1	0.50	0.42	0.72	0.63	0.58	0.56	0.79	0.83	0.99	0.70
Approach 2	0.45	0.39	0.61	0.53	0.56	0.41	0.75	0.62	0.89	0.59
Approach 3	0.48	0.42	0.66	0.57	0.60	0.45	0.76	0.65	0.91	0.65
Approach 4	0.52	0.57	0.74	0.58	0.61	0.58	0.76	0.67	1.00	0.69
Approach 5	0.65	0.46	0.79	0.69	0.81	0.64	0.89	0.72	1.00	0.76
Approach 6	0.56	0.40	0.67	0.61	0.67	0.50	0.77	0.70	0.94	0.60
Approach 7	0.48	0.40	0.63	0.54	0.58	0.43	0.77	0.63	0.90	0.60
Approach 8	0.51	0.53	0.66	0.57	0.62	0.46	0.88	0.73	0.93	0.61
Approach 9	0.61	0.57	0.77	0.54	0.70	0.59	0.94	0.75	1.00	0.62
Approach 10	0.50	0.48	0.64	0.53	0.63	0.47	0.84	0.64	0.95	0.67

30

## Benchmarks - Ranks

	Data 1	Data 2	Data 3	Data 4	Data 5	Data 6	Data 7	Data 8	Data 9	Data 10
Approach 1	6	7	4	2	8	4	5	1	4	2
Approach 2	10	10	10	10	10	10	10	10	10	10
Approach 3	8	6	6	5	7	8	9	7	8	5
Approach 4	4	1	3	4	6	3	8	6	1	3
Approach 5	1	5	1	1	1	1	2	4	1	1
Approach 6	3	8	5	3	3	5	6	5	6	9
Approach 7	9	9	9	7	9	9	7	9	9	8
Approach 8	5	3	7	6	5	7	3	3	7	7
Approach 9	2	2	2	8	2	2	1	2	1	6
Approach 10	7	4	8	9	4	6	4	8	5	4

31

## Benchmarks - Average Ranks

	Data 1	Data 2	Data 3	Data 4	Data 5	Data 6	Data 7	Data 8	Data 9	Data 10	Avg Rank
Approach 1	6	7	4	2	8	4	5	1	4	2	4.3
Approach 2	10	10	10	10	10	10	10	10	10	10	10.0
Approach 3	8	6	6	5	7	8	9	7	8	5	6.9
Approach 4	4	1	3	4	6	3	8	6	1	3	3.9
Approach 5	1	5	1	1	1	1	2	4	1	1	1.8
Approach 6	3	8	5	3	3	5	6	5	6	9	5.3
Approach 7	9	9	9	7	9	9	7	9	9	8	8.5
Approach 8	5	3	7	6	5	7	3	3	7	7	5.3
Approach 9	2	2	2	8	2	2	1	2	1	6	2.8
Approach 10	7	4	8	9	4	6	4	8	5	4	5.9

32



## Benchmarks - Average Ranks

	Data 1	Data 2	Data 3	Data 4	Data 5	Data 6	Data 7	Data 8	Data 9	Data 10	Avg Rank
Approach 1	6	7	4	2	8	4	5	1	4	2	4.3
Approach 2	10	10	10	10	10	10	10	10	10	10	10.0
Approach 3	8	6	6	5	7	8	9	7	8	5	6.9
Approach 4	4	1	3	4	6	3	8	6	1	3	3.9
<b>Approach 5</b>	<b>1</b>	<b>5</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>2</b>	<b>4</b>	<b>1</b>	<b>1</b>	<b>1.8</b>
Approach 6	3	8	5	3	3	5	6	5	6	9	5.3
Approach 7	9	9	9	7	9	9	7	9	9	8	8.5
Approach 8	5	3	7	6	5	7	3	3	7	7	5.3
Approach 9	2	2	2	8	2	2	1	2	1	6	2.8
Approach 10	7	4	8	9	4	6	4	8	5	4	5.9

33

## Benchmarks - Significance Testing

Another characteristic of academic evaluations versus industry evaluations is the use of significance testing

– Is the *best* method really *best*?

34

## Benchmarks - Significance Testing

We recommend a two step process:

- **Friedman aligned rank test** to first test whether a significant difference between the performance of the algorithms over the datasets exists
  - Compared against a significance level (e.g. 0.05)
- If a difference does exist then a pairwise **Nemenyi test** should be performed to show between which algorithm pairs the significant differences exist
  - Compared against a significance level (e.g. 0.05)

35

## Benchmarks - Significance Testing

Nemenyi test gives us a significance matrix

- Which approaches are significantly different from which others?

Significance matrix can give rise to a **critical differences plot** which shows groups of approaches which are significantly different from each other

36

## Benchmarks - Significance Testing

### Friedman's Aligned Rank Test for Multiple Comparisons

–  $T = 57.609$ ,  $df = 9$ ,  $p\text{-value} = 3.862e-09$

### Nemenyi test

– Critical difference = 4.393,  $k = 10$ ,  $df = 90$

37

## Benchmarks - Significance Testing

	M1	M2	M3	M4	M5	M6	M7	M8	M9	M10
M1	0	-5.6	-2.7	0.25	2.45	-1	-4.05	-0.85	1.55	-1.55
M2	-5.6	0	2.9	5.85	8.05	4.6	1.55	4.75	7.15	4.05
M3	-2.7	2.9	0	2.95	5.15	1.7	-1.35	1.85	4.25	1.15
M4	0.25	5.85	2.95	0	2.2	-1.25	-4.3	-1.1	1.3	-1.8
M5	2.45	8.05	5.15	2.2	0	-3.45	-6.5	-3.3	-0.9	-4
M6	-1	4.6	1.7	-1.25	-3.45	0	-3.05	0.15	2.55	-0.55
M7	-4.05	1.55	-1.35	-4.3	-6.5	-3.05	0	3.2	5.6	2.5
M8	-0.85	4.75	1.85	-1.1	-3.3	0.15	3.2	0	2.4	-0.7
M9	1.55	7.15	4.25	1.3	-0.9	2.55	5.6	2.4	0	-3.1
M10	-1.55	4.05	1.15	-1.8	-4	-0.55	2.5	-0.7	-3.1	0

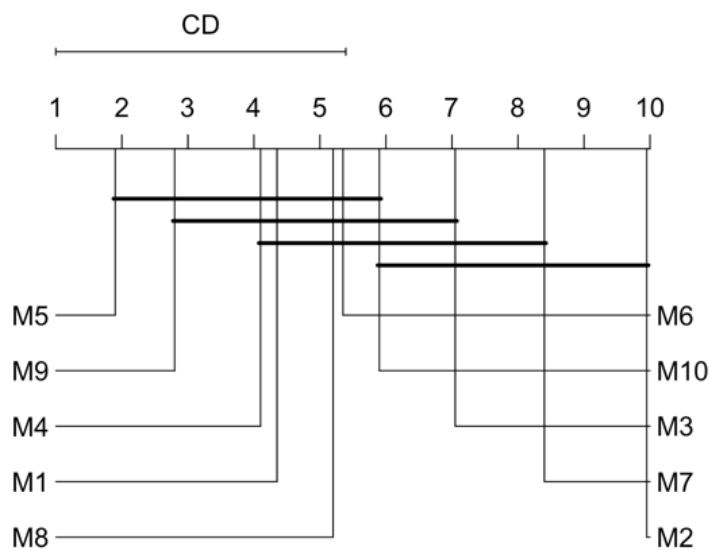
38

## Benchmarks - Significance Testing

	M1	M2	M3	M4	M5	M6	M7	M8	M9	M10
M1	False	True	False	False	False	False	False	False	False	False
M2	True	False	False	True	True	True	False	True	True	False
M3	False	False	False	False	True	False	False	False	False	False
M4	False	True	False	False	False	False	False	False	False	False
M5	False	True	True	False	False	False	True	False	False	False
M6	False	True	False	False	False	False	False	False	False	False
M7	False	False	False	False	True	False	False	False	True	False
M8	False	True	False	False	False	False	False	False	False	False
M9	False	True	False	False	False	False	True	False	False	False
M10	False	False	False	False	False	False	False	False	False	False

39

## Benchmarks - Significance Testing



Critical  
differences  
plot

40

# The Great Time Series Bakeoff


Data Min Knowl Disc (2017) 31:606–660  
DOI 10.1007/s10618-016-0483-9



**Table 6** Average accuracy of the best nine classifiers over 85 problems

Datasets	COTE	ST	BOSS	EE	DTW <sub>F</sub>	TSF	TSBF	LPS	MSM
Adiac	<b>0.81</b>	0.768	0.749	0.665	0.605	0.707	0.727	0.765	
ArrowHead	<b>0.877</b>	0.851	0.875	0.86	0.776	0.789	0.801	0.806	
Beef	<b>0.764</b>	0.736	0.615	0.532	0.546	0.648	0.554	0.52	
BeetleFly	0.921	0.875	<b>0.949</b>	0.823	0.853	0.842	0.799	0.893	
BirdChicken	0.941	0.927	<b>0.984</b>	0.848	0.865	0.839	0.902	0.854	
Car	0.899	<b>0.902</b>	0.855	0.799	0.851	0.758	0.795	0.836	
CBF	0.998	0.986	<b>0.998</b>	0.993	0.979	0.958	0.977	0.984	
ChlorineConcentration	<b>0.736</b>	0.682	0.66	0.659	0.658	0.719	0.683	0.642	
CinCECGtorso	<b>0.983</b>	0.918	0.9	0.946	0.714	0.974	0.716	0.743	
Coffee	<b>1</b>	0.995	0.989	0.989	0.973	0.989	0.982	0.95	
Computers	0.77	0.785	<b>0.802</b>	0.732	0.659	0.768	0.765	0.726	
CricketX	<b>0.814</b>	0.777	0.764	0.801	0.769	0.691	0.731	0.696	
CricketY	<b>0.815</b>	0.762	0.749	0.794	0.756	0.688	0.728	0.706	
CricketZ	<b>0.827</b>	0.798	0.776	0.804	0.785	0.707	0.738	0.714	
DiatomSizeReduction	0.925	0.911	0.939	<b>0.946</b>	0.942	0.941	0.89	0.915	
DistalPhalanxOAG	0.805	<b>0.829</b>	0.815	0.768	0.796	0.809	0.816	0.767	
DistalPhalanxOC	<b>0.821</b>	0.819	0.814	0.768	0.76	0.813	0.812	0.742	
DistalPhalanxTW	<b>0.693</b>	0.69	0.673	0.654	0.658	0.686	0.69	0.618	
Earthquakes	0.747	0.737	0.746	0.735	<b>0.747</b>	0.747	0.747	0.668	
ECG200	0.873	0.84	<b>0.89</b>	0.881	0.819	0.868	0.847	0.807	

classification bake off: a review  
evaluation of recent algorithmic



The great time series classification bake off: a review and experimental evaluation of recent algorithmic advances, Bagnall et al

<https://link.springer.com/article/10.1007/S10618-016-0483-9>

Time Series Classification

<https://www.timeseriesclassification.com/results.php>

41

## Summary

42


## Summary

### Evaluation

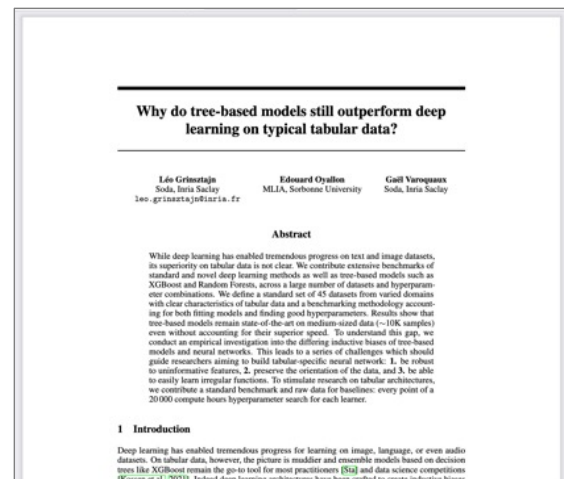
- Choosing appropriate evaluation mechanisms is crucial in doing machine learning properly
- Macro versus micro averaging is a mistake too often made
- One of the key differences between *industry* evaluations and *research* evaluations is the need for significance testing
- Lots of research evaluations reduce to a benchmark across multiple methods on multiple datasets

43

## Discussion



"Common pitfalls and recommendations for using machine learning to detect and prognosticate for COVID-19 ...", Roberts et al  
<https://www.nature.com/articles/s42256-021-00307-0>



"Why do tree-based models still outperform deep learning on typical tabular data?", Grinsztajn et al  
[https://openreview.net/pdf?id=Fp7\\_phQsxn](https://openreview.net/pdf?id=Fp7_phQsxn)

44

## Questions

